

Adding AD to existing projects

SIAM CSE 2021



U.S. DEPARTMENT OF
ENERGY

Office of
Science

What we usually start with

- ▶ Existing code spread over many files
- ▶ Non trivial build system
- ▶ Multiple build time configuration options
- ▶ Multiple top level routines and drivers
- ▶ Calls to libraries
- ▶ Extensive I/O, initialization
- ▶ Parallelism
- ▶ Many users and use cases not needing AD

Code Preparation

Analyze the code . . .

- ▶ Start with a small case first
- ▶ Identify portion(s) and top level function(s) that need to be differentiated
- ▶ Identify independent and dependent variables
- ▶ Identify initialization code and I/O
- ▶ Identify portions that cannot be differentiated

Code Modification and Maintenance

Prepare the code . . .

- ▶ Rewrite portions the cannot be differentiated
 - ▶ Special handling for solvers
 - ▶ Stubs for library calls
 - ▶ Code changes may need to *guarded* using preprocessor directives
- ▶ Write preprocessing scripts
- ▶ Add AD to existing build system
- ▶ Regression tests to ensure that changes for AD do not change primal output

Validation and Debug


- ▶ Original Primal vs. AD Primal
- ▶ Finite Differences vs Forward mode¹
- ▶ Forward mode vs Reverse mode
- ▶ Many tools provide feedback on problems.
- ▶ Start always with the simplest case possible

¹See https://fluids.ac.uk/files/resources/Farrell_adjoint_slides.pdf p.6 for a more accurate method

The dot-product test

The dot-product of all active tangent variables with all active adjoint variables remains the same throughout the code.

$$y = h(g(f(x)))$$



Forward: \dot{x}


Reverse: $\partial f \cdot \partial g \cdot \partial h \cdot \bar{y}$

Dot-product: $\dot{x} \cdot \partial f \cdot \partial g \cdot \partial h \cdot \bar{y}$

The dot-product test

The dot-product of all active tangent variables with all active adjoint variables remains the same throughout the code.

$$y = h(g(f(x)))$$



Forward: $\dot{x} \cdot \partial f$


Reverse: $\partial g \cdot \partial h \cdot \bar{y}$

Dot-product: $\dot{x} \cdot \partial f \cdot \partial g \cdot \partial h \cdot \bar{y}$

The dot-product test

The dot-product of all active tangent variables with all active adjoint variables remains the same throughout the code.

$$y = h(g(f(x)))$$



Forward: $\dot{x} \cdot \partial f \cdot \partial g$


Reverse: $\partial h \cdot \bar{y}$

Dot-product: $\dot{x} \cdot \partial f \cdot \partial g \cdot \partial h \cdot \bar{y}$

The dot-product test

The dot-product of all active tangent variables with all active adjoint variables remains the same throughout the code.

$$y = h(g(f(x)))$$



Forward: $\dot{x} \cdot \partial f \cdot \partial g \cdot \partial h$

Reverse: \bar{y}

Dot-product: $\dot{x} \cdot \partial f \cdot \partial g \cdot \partial h \cdot \bar{y}$

Limitations of the dot-product test

1. **It only checks for a given seed.** The dot-product may look good if some errors cancel each other out. (Given some correct tangent-linear vector and a dot-product result, there's an infinite number of adjoint vectors with the same dot-product result).
 - ▶ Dot-product is a necessary, but not a sufficient condition for correctness (except for scalar functions). Sufficient condition for consistency: Dot-product test holds for a complete set of basis vectors.
2. **Results will differ by some small amount.** Is it roundoff, or a bug?
3. **It only checks for a given primal input.** If the dot-product works at some point, it may not work somewhere else.
4. It only finds discrepancies between forward and reverse, assuming that forward is correct