

# Memory Requirements

SIAM CSE 2021



U.S. DEPARTMENT OF  
**ENERGY**

Office of  
Science

# Memory requirements of AD

Computing derivatives increases the program memory footprint. We must store:

- ▶ the intermediate and final derivatives
- ▶ (*depending on AD mode*) intermediate values, partial derivatives, sequence of primal operations, trace of control flow. . .

Alternatives to AD (continuous adjoints, FD) face similar issues.

A challenge especially for reverse AD (cheap gradients, but “no free lunch”).

# Overloading AD

Overloading AD is traditional name for **taping** a trace of the full forward run, e.g. for each *run-time* statement  $x := a \text{ Op } b$

- ▶ [long int]: an index representing  $a$  (and  $a'$ ),
- ▶ [long int]: an index representing  $b$  (and  $b'$ ),
- ▶ [short int]: a code for  $\text{Op}$ ,
- ▶ [double]: the value of  $x$  *before* the statement.

Tape grows linearly with execution time.

Tape contains everything  $\Rightarrow$  no backward sweep code, only a conceptually simple tape interpreter.

Overloading AD is probably the most memory-consuming AD.

# Source-Transformation AD

ST-AD generates a new source program  $P'$ , replacing a good part of the tape with source in  $P'$ .

- ▶ No indices stored, nor op-codes:  
only new variables  $a'$ ,  $b'$ ,  $x'$
- ▶ Stack memory footprint of  $P$  (only) doubled
- ▶ End-user *and* compiler can take a look at  $P'$

Reverse AD still tapes values for data-flow reversal

⇒ tape is smaller, still grows linear with run-time.

What can we do?

# Memory requirements of Data-flow reversal

Recall the concept of Data-flow reversal:




Assume **red-arrow magic** done by storing on a stack

Step back a little:




Stack grows like run-time, reaching maximum at turn point (on the right).

# Checkpointing: trading recomputation for storage



Checkpointing C is:




- ▶ **reduces** peak storage
- ▶ at the cost of **duplicate execution**
- ▶ also costs a memory “**Snapshot**”, small enough:

$$\text{Snapshot} \subset \text{use}(\overline{C}) \cap (\text{out}(C) \cup \text{out}(\overline{D}))$$

# Nesting Checkpoints

On large codes, checkpoints can/must be nested.



On a well-balanced nesting (e.g. a well-balanced call tree), memory and CPU grow like  $\log(\text{runtime})$

## Technical constraints

- ▶ start/end of C in the same procedure
- ▶ start/end of C in the same control

Also, C must be **reentrant**, i.e. one can restore exactly its initial state:


- ▶ if C contains a `malloc`, it must contain its `free`
- ▶ if C contains a `send`, it must contain its `recv`
- ▶ if C contains a `isend/irecv`, it must contain its `wait`
- ▶ if C contains a `open`, it must contain its `close`

Most often C's are **procedures, time steps** ...



# Checkpointing on Time-stepping loops


**Binomial** checkpointing nesting is optimal for time-stepping loops (uniform step cost, depend on previous step only, negligible snapshot time):



- ▶ peak memory storage and duplicate execution still grow like  $\log(\text{runtime})$ , but proved optimal.
- ▶ in real life, storage is fixed to  $q$  snapshots, execution duplication grows like  $\sqrt[q]{\text{runtime}}$

# Checkpointing on the Call Tree

The **Call Tree** is the natural support to define your checkpointing strategy:



- Costs still grow like  **$\log(\text{runtime})$**  if call tree well balanced.
- Ill-balanced call trees require **not** checkpointing some calls.
- Small leaf procedures better **not** checkpointed.

# Checkpointing on the Call Tree

foo not checkpointed (aka <b>Split</b> )	foo checkpointed (aka <b>Joint</b> )
<pre>... <math>\overrightarrow{\hspace{1cm}}</math>   foo(a) ... <math>\overleftarrow{\hspace{1cm}}</math>   foo(a, <math>\bar{a}</math>) ... <math>\overrightarrow{\hspace{1cm}}</math> foo(x)   push x   x = sin(x)   push x   x = x*x } <math>\overleftarrow{\hspace{1cm}}</math> foo(x, <math>\bar{x}</math>)   pop x   <math>\bar{x} = 2*x*\bar{x}</math>   pop x   <math>\bar{x} = \cos(x)*\bar{x}</math> }</pre>	<pre>...   push a   foo(a) ...   pop a   <math>\overleftarrow{\hspace{1cm}}</math>   foo(a, <math>\bar{a}</math>) ...    <math>\overleftarrow{\hspace{1cm}}</math>   foo(x, <math>\bar{x}</math>)     push x     x = sin(x)     <math>\bar{x} = 2*x*\bar{x}</math>     pop x     <math>\bar{x} = \cos(x)*\bar{x}</math>   }</pre>

# Advanced research on Data-flow reversal

- ▶ combine storage of intermediates/partials, and inversion / forward recalculation.
- ▶ periodic checkpointing
- ▶ binary checkpointing
- ▶ optimal strategies if non-zero snapshot time
- ▶ combination with resilience checkpoints